

Apertus baseline scaling — Apertus-8B (left half) vs Apertus-70B (right half) per bucket (no-skill, python-tiny)  
solved % above ( $>0.8$ ); failed % below ( $\leq 0.05$ ); colour-coded median line per model

