

Apertus baseline scaling — Apertus-8B (left half) vs Apertus-70B (right half) per bucket (no-skill, python-tiny)
solved % above (>0.8); failed % below (≤ 0.05); colour-coded median line per model

