

Qwen3 baseline scaling — Qwen3-8B (left half) vs Qwen3-32B (right half) per bucket (no-skill, python-tiny)
solved % above (>0.8); failed % below (≤ 0.05); colour-coded median line per model

