

RQ1: does skill-v2.1-sys help vs no-skill, per model? max(edit, winnowing), python-tiny
left half = baseline; right half = skill-v2.1-sys; solved % above (>0.8); failed % below (≤ 0.05)

